

COMP9024: Data Structures and Algorithms




Sorting Algorithms



Outline

- Merge Sort
- Quick Sort
- Bucket-Sort
- Radix Sort
- Sorting Lower Bound

Merge Sort





Divide-and-Conquer

- Divide-and conquer is a general algorithm design paradigm:
 - **Divide**: divide the input data S in two disjoint subsets S_1 and S_2
 - **Recur**: solve the subproblems associated with S_1 and S_2
 - **Conquer**: combine the solutions for S_1 and S_2 into a solution for S
- The base case for the recursion are subproblems of size 0 or 1
- Merge-sort is a sorting algorithm based on the divide-and-conquer paradigm
- Like heap-sort
 - It has $O(n \log n)$ running time
- Unlike heap-sort
 - It accesses data in a sequential manner (suitable to sort data on a disk)

Merge-Sort

- Merge-sort on an input sequence S with n elements consists of three steps:
 - **Divide**: partition S into two sequences S_1 and S_2 of about $n/2$ elements each
 - **Recur**: recursively sort S_1 and S_2
 - **Conquer**: merge S_1 and S_2 into a unique sorted sequence

Algorithm *mergeSort*(S)

Input sequence S with n elements

Output sorted sequence S

```
{  
  if (  $S.size() > 1$  )  
  {  
     $(S_1, S_2) = partition(S, n/2);$   
    mergeSort( $S_1$ );  
    mergeSort( $S_2$ );  
     $S = merge(S_1, S_2);$   
  }  
}
```

Merging Two Sorted Sequences

- The conquer step of merge-sort consists of merging two sorted sequences A and B into a sorted sequence S containing the union of the elements of A and B
- Merging two sorted sequences, each with $n/2$ elements and implemented by means of a doubly linked list, takes $O(n)$ time

Algorithm *merge*(A, B)


Input sequences A and B with $n/2$ elements each

Output sorted sequence of $A \cup B$

```
{  
   $S$  = empty sequence;  
  while (  $\neg isEmpty(A) \wedge \neg isEmpty(B)$  )  
    if (  $firstElement(A) < firstElement(B)$  )  
       $insertLast(S, removeFirst(A))$ ;  
    else  
       $insertLast(S, removeFirst(B))$ ;  
  while (  $\neg isEmpty(A)$  )  
     $insertLast(S, removeFirst(A))$ ;  
  while (  $\neg isEmpty(B)$  )  
     $insertLast(S, removeFirst(B))$ ;  
  return  $S$ ;  
}
```


Merge-Sort Tree

- An execution of merge-sort is depicted by a binary tree
 - each node represents a recursive call of merge-sort and stores
 - unsorted sequence before the execution and its partition
 - sorted sequence at the end of the execution
 - the root is the initial call
 - the leaves are calls on subsequences of size 0 or 1




Execution Example (1/10)

- Partition




Execution Example (2/10)

- Recursive call, partition




Execution Example (3/10)

- Recursive call, partition




Execution Example (4/10)

- Recursive call, base case




Execution Example (5/10)

- Recursive call, base case



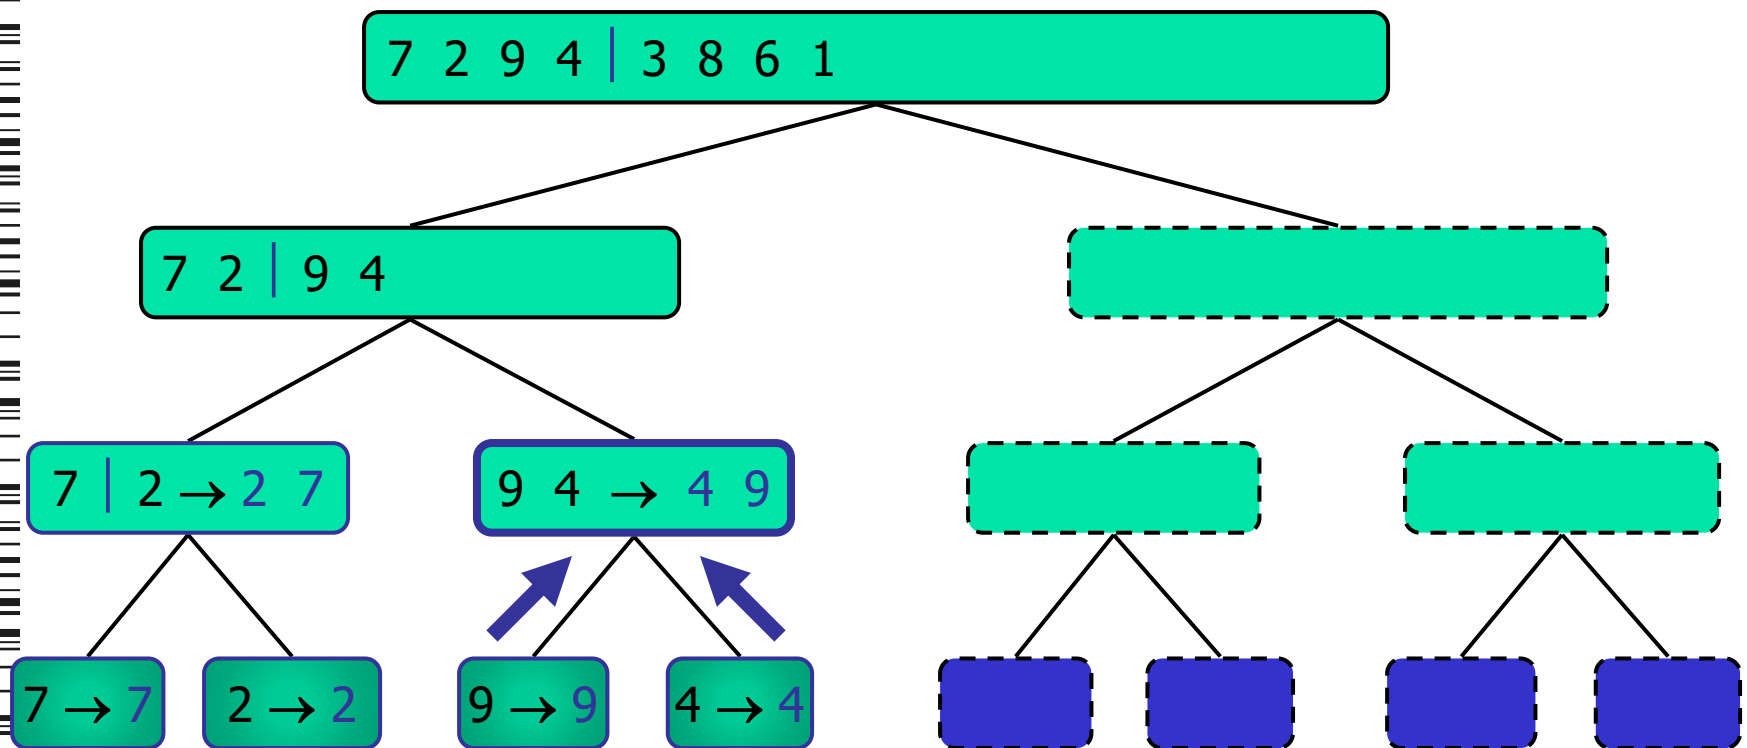
Execution Example (6/10)

- Merge



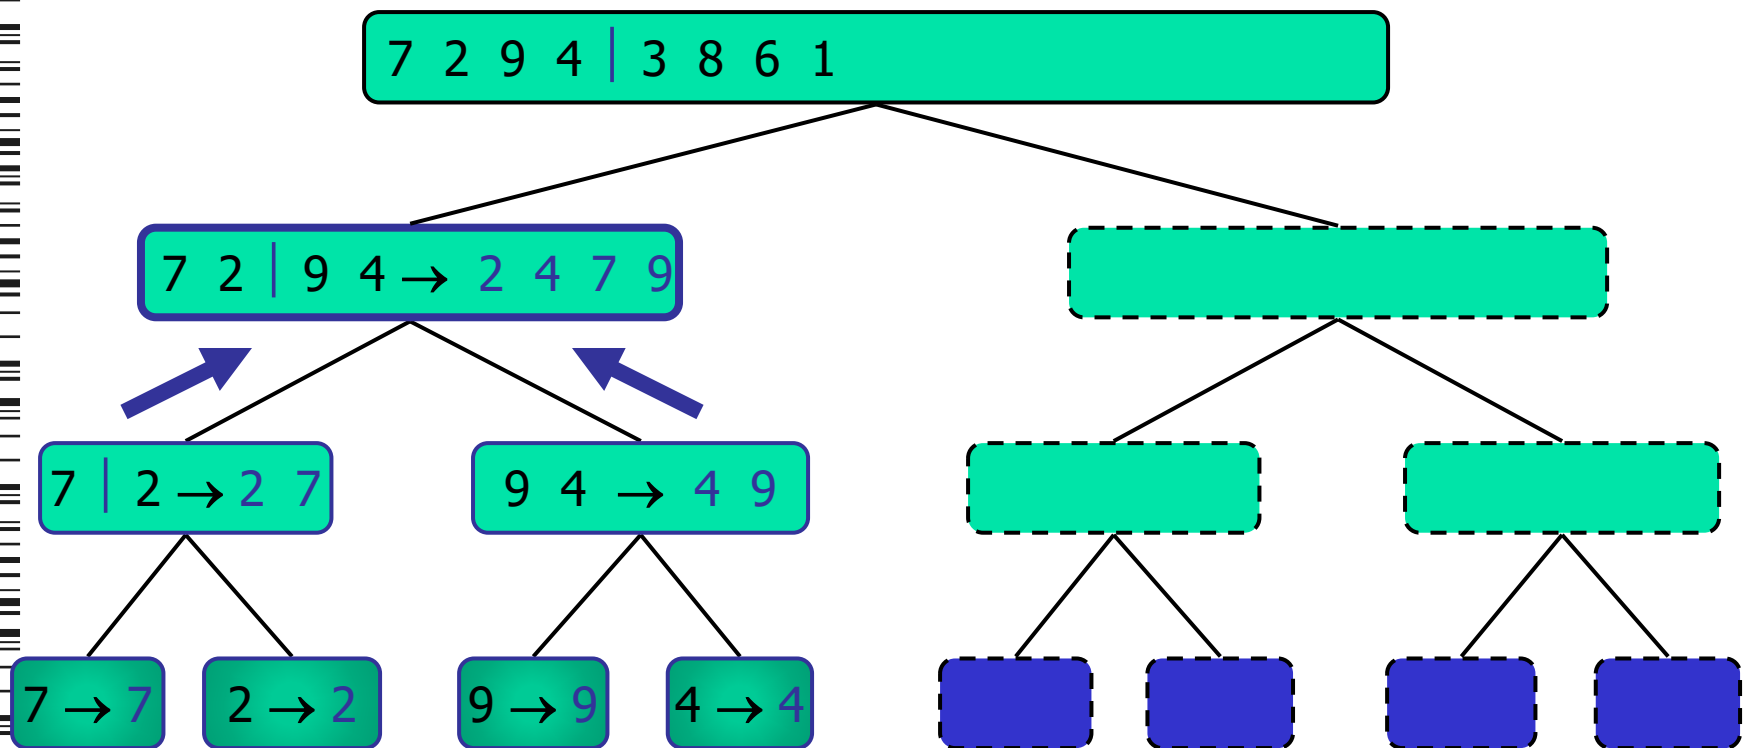
Execution Example (7/10)

- Recursive call, ..., base case, merge



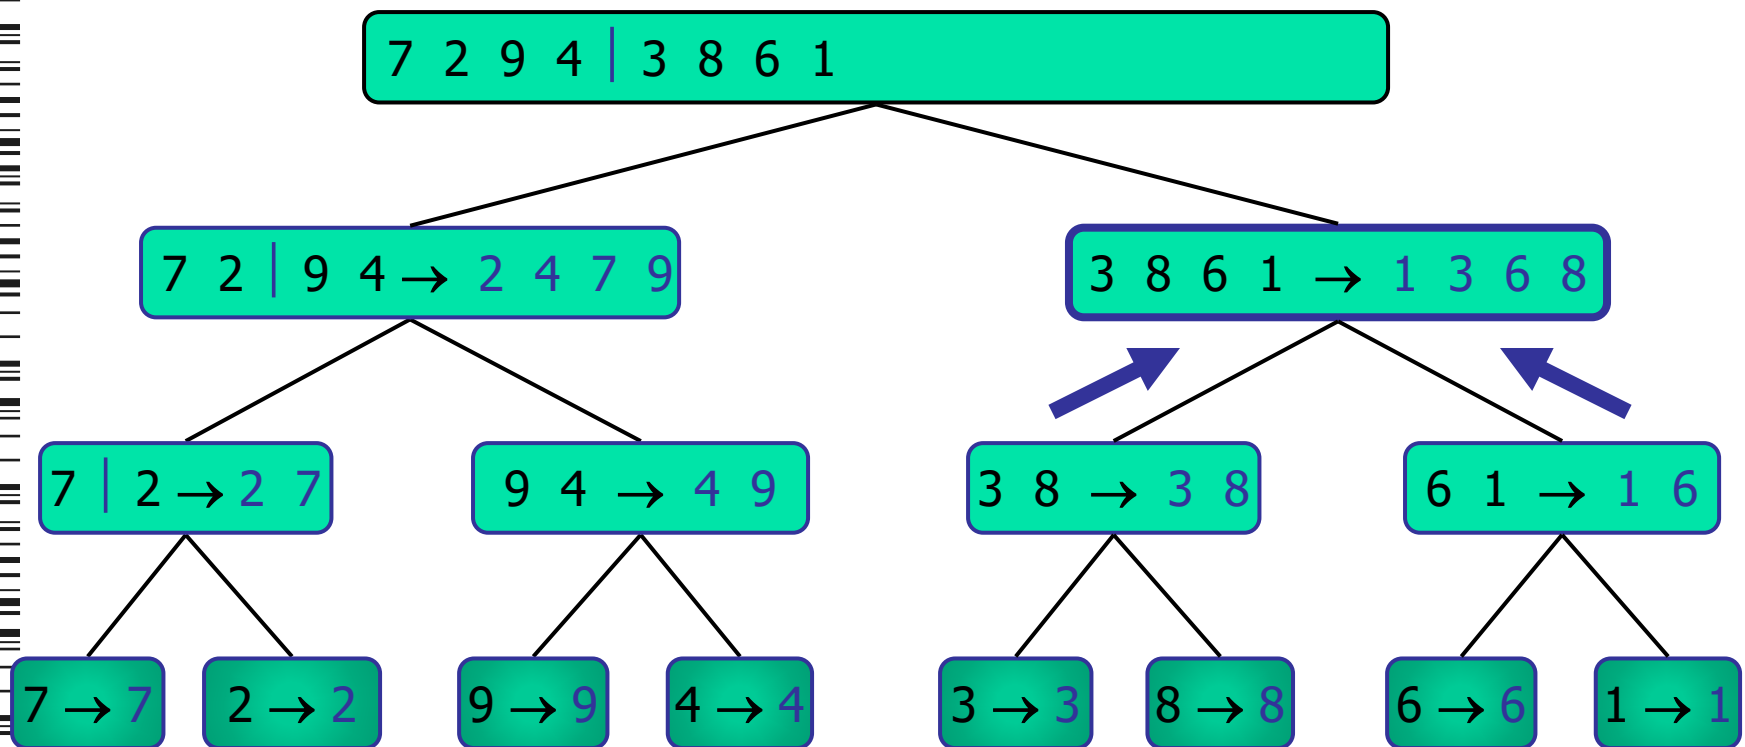
Execution Example (8/10)

- Merge




Execution Example (9/10)

- Recursive call, ..., merge, merge



Execution Example (10/10)

- Merge



Analysis of Merge-Sort

- The height h of the merge-sort tree is $O(\log n)$
 - at each recursive call we halve the sequence,
- The overall amount of work done at the nodes of depth i is $O(n)$
 - we partition and merge 2^i sequences of size $n/2^i$
 - we make 2^{i+1} recursive calls
- Thus, the total running time of merge-sort is $O(n \log n)$


depth #seqs size

0 1 n

1 2 $n/2$

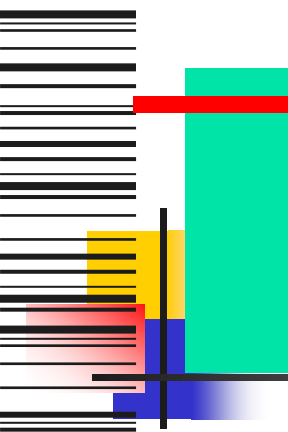
i 2^i $n/2^i$

... ...




Summary of Sorting Algorithms

Algorithm	Time	Notes
selection-sort	$O(n^2)$	<ul style="list-style-type: none">■ slow■ in-place■ for small data sets (< 1K)
insertion-sort	$O(n^2)$	<ul style="list-style-type: none">■ slow■ in-place■ for small data sets (< 1K)
heap-sort	$O(n \log n)$	<ul style="list-style-type: none">■ fast■ in-place■ for large data sets (1K — 1M)
merge-sort	$O(n \log n)$	<ul style="list-style-type: none">■ fast■ sequential data access■ for huge data sets (> 1M)



Recursive Merge-Sort on Array (1/3)



// Array A[] is the source array to sort; array B[] is a temporary array.

```
void TopDownMergeSort(A[], B[], n)
```

```
{
```

```
    CopyArray(A, 0, n, B);    // one time copy of A[] to B[]
```

```
    TopDownSplitMerge(B, 0, n, A); // sort data from B[] into A[]
```

```
}
```

```
void CopyArray(A[], iBegin, iEnd, B[])
```

```
{
```

```
    for(k = iBegin; k < iEnd; k++)
```

```
        B[k] = A[k];
```

```
}
```

Recursive Merge-Sort on Array (2/3)

```
// Sort the given run of array A[] using array B[] as a source.
// iBegin is inclusive; iEnd is exclusive (A[iEnd] is not in the set).
void TopDownSplitMerge(B[], iBegin, iEnd, A[])
{
    if(iEnd - iBegin < 2)           // if run size == 1
        return;                   // consider it sorted
    // split the run longer than 1 item into halves
    iMiddle = (iEnd + iBegin) / 2;  // iMiddle = mid point
    // recursively sort both runs from array A[] into B[]
    TopDownSplitMerge(A, iBegin, iMiddle, B); // sort the left run
    TopDownSplitMerge(A, iMiddle, iEnd, B); // sort the right run
    // merge the resulting runs from array B[] into A[]
    TopDownMerge(B, iBegin, iMiddle, iEnd, A);
}
```

Recursive Merge-Sort on Array (3/3)

```
// Left source half is A[ iBegin:iMiddle-1].
// Right source half is A[iMiddle:iEnd-1  ].
// Result is B[iBegin:iEnd-1].
void TopDownMerge(A[], iBegin, iMiddle, iEnd, B[])
{ i = iBegin, j = iMiddle;
  // while there are elements in the left or right runs
  for (k = iBegin; k < iEnd; k++) {
    // If left run head exists and is <= existing right run head.
    if (i < iMiddle && (j >= iEnd || A[i] <= A[j])) {
      B[k] = A[i];
      i = i + 1;
    } else {
      B[k] = A[j];
      j = j + 1;
    }
  }
}
```


Nonrecursive Merge-Sort on Array (1/2)

```
// array A[] has the items to sort; array B[] is a work array
void BottomUpMergeSort(A[], B[], n)
{ // Each 1-element run in A is already "sorted".
  // Make successively longer sorted runs of length 2, 4, 8, 16... until whole array is sorted.
  for (width = 1; width < n; width = 2 * width)
  { // Array A is full of runs of length width.
    for (i = 0; i < n; i = i + 2 * width)
    { // Merge two runs: A[i:i+width-1] and A[i+width:i+2*width-1] to B[]
      // or copy A[i:n-1] to B[] ( if(i+width >= n) )
      BottomUpMerge(A, i, min(i+width, n), min(i+2*width, n), B);
    }
    // Now work array B is full of runs of length 2*width.
    // Copy array B to array A for next iteration.
    // A more efficient implementation would swap the roles of A and B.
    CopyArray(B, 0, n, A);
    // Now array A is full of runs of length 2*width.
  }
}
```

Nonrecursive Merge-Sort on Array (2/2)

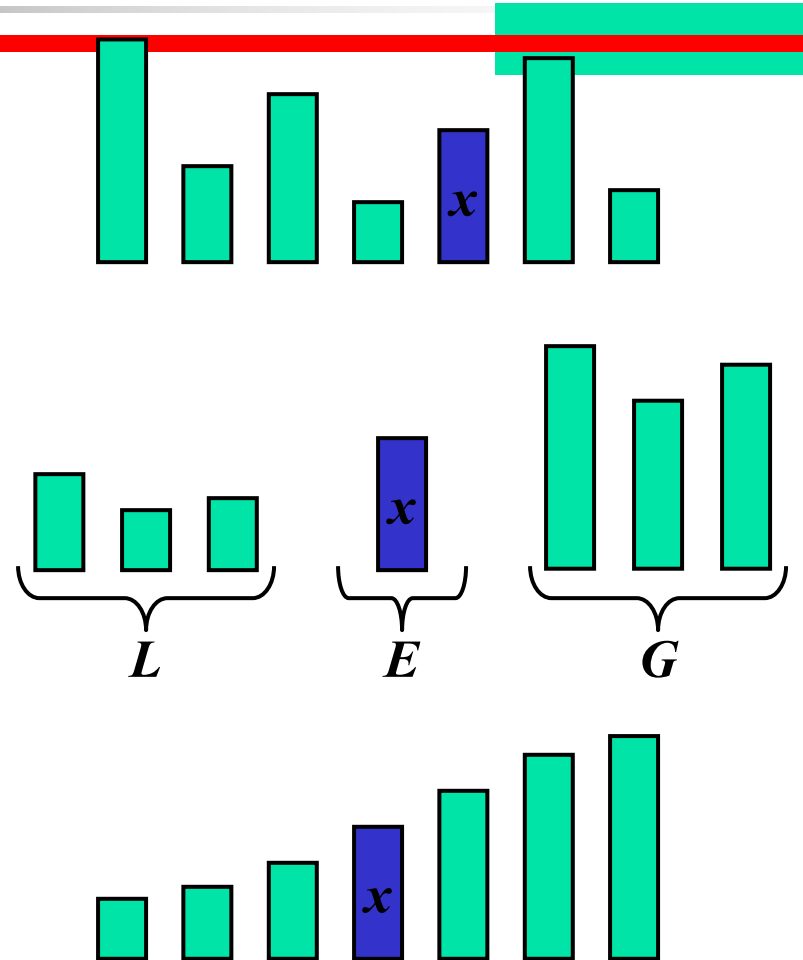
```
// Left run is A[iLeft :iRight-1].
// Right run is A[iRight:iEnd-1 ].
void BottomUpMerge(A[], iLeft, iRight, iEnd, B[])
{
    i = iLeft, j = iRight;
    // While there are elements in the left or right runs...
    for (k = iLeft; k < iEnd; k++) {
        // If left run head exists and is <= existing right run head.
        if (i < iRight && (j >= iEnd || A[i] <= A[j])) {
            B[k] = A[i];
            i = i + 1;
        } else {
            B[k] = A[j];
            j = j + 1;
        }
    }
}
```


Quick-Sort




Quick-Sort

- Quick-sort is a randomized sorting algorithm based on the divide-and-conquer paradigm:
 - **Divide**: pick a random element x (called **pivot**) and partition S into
 - L elements less than x
 - E elements equal x
 - G elements greater than x
 - **Recur**: sort L and G
 - **Conquer**: join L , E and G



Partition



- We partition an input sequence as follows:
 - We remove, in turn, each element y from S and
 - We insert y into L , E or G , depending on the result of the comparison with the pivot x
- Each insertion and removal is at the beginning or at the end of a sequence, and hence takes $O(1)$ time
- Thus, the partition step of quick-sort takes $O(n)$ time

Algorithm *partition*(S, p)


Input sequence S , position p of pivot

Output subsequences L , E , G of the elements of S less than, equal to, or greater than the pivot, resp.

```
{  $L, E, G$  = empty sequences;  
   $x = S.remove(p)$ ;  
  while (  $\neg S.isEmpty()$  )  
    {  $y = S.remove(S.first())$ ;  
      if (  $y < x$  )  
         $L.insertLast(y)$ ;  
      else if (  $y = x$  )  
         $E.insertLast(y)$ ;  
      else //  $y > x$   
         $G.insertLast(y)$ ;  
    }  
  return  $L, E, G$ ;  
}
```


Quick-Sort Tree

- An execution of quick-sort is depicted by a binary tree
 - Each node represents a recursive call of quick-sort and stores
 - Unsorted sequence before the execution and its pivot
 - Sorted sequence at the end of the execution
 - The root is the initial call
 - The leaves are calls on subsequences of size 0 or 1




Execution Example (1/7)

- Pivot selection




Execution Example (2/7)

- Partition, recursive call, pivot selection




Execution Example (3/7)

- Partition, recursive call, base case



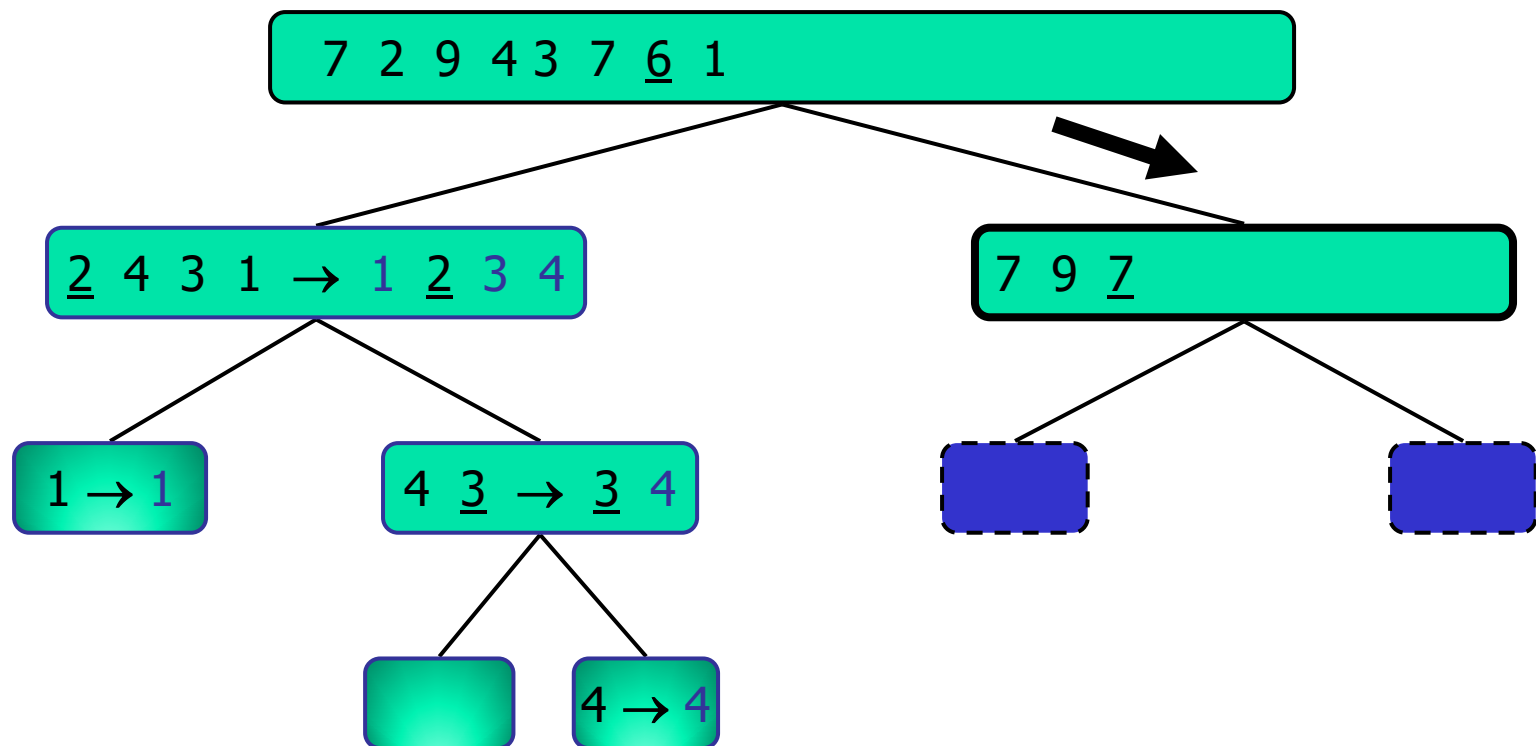
Execution Example (4/7)

- Recursive call, ..., base case, join




Execution Example (5/7)

- Recursive call, pivot selection




Execution Example (6/7)

- Partition, ..., recursive call, base case



Execution Example (7/7)

- Join, join



Worst-case Running Time

- The worst case for quick-sort occurs when the pivot is the unique minimum or maximum element
- One of L and G has size $n - 1$ and the other has size 0
- The running time is proportional to the sum

$$n + (n - 1) + \dots + 2 + 1$$

- Thus, the worst-case running time of quick-sort is $O(n^2)$


depth time

0 n

1 $n - 1$


...

$n - 1$ 1




Expected Running Time (1/2)

- Consider a recursive call of quick-sort on a sequence of size s
 - **Good call:** the sizes of L and G are each less than $3s/4$
 - **Bad call:** one of L and G has size greater than $3s/4$



Good call




Bad call

- A call is **good** with probability $1/2$
 - $1/2$ of the possible pivots cause good calls:



Expected Running Time (2/2)

- **Probabilistic Fact:** The expected number of coin tosses required in order to get k heads is $2k$
- For a node of depth i , we expect
 - $i/2$ ancestors are good calls
 - The size of the input sequence for the current call is at most $(3/4)^{i/2}n$
- Therefore, we have
 - For a node of depth $2\log_{4/3}n$, the expected input size is one
 - The expected height of the quick-sort tree is $O(\log n)$
- The amount of work done at the nodes of the same depth is $O(n)$
- Thus, the expected running time of quick-sort is $O(n \log n)$



total expected time: $O(n \log n)$

In-Place Quick-Sort



- Quick-sort can be implemented to run in-place
- In the partition step, we rearrange the elements of the input sequence such that
 - the elements less than the pivot have indices less than p
 - the elements equal to or greater than the pivot have indices between $p+1$ and r
- The recursive calls consider
 - elements with indices less than p .
 - elements with indices greater than p

Algorithm *inPlaceQuickSort*(S, l, r)

Input sequence S , indices l and r of the first and last elements

Output sequence S with the elements of indices between l and r rearranged in increasing order

```
{ if  $l \geq r$ 
    return;
 $k$  = a random integer between  $l$  and  $r$ ;
pivot =  $S[k]$ ;
Swap pivot and  $S[r]$ ;
// pivot is the last element now
 $p$  = inPlacePartition();
inPlaceQuickSort( $S, l, p-1$ );
inPlaceQuickSort( $S, p+1, r$ );
}
```

In-Place Partitioning (1/3)




```
inPlacePartition()
```

```
{ left = l;           // scan right to locate the end of G U E
  right = r;          // scan left to locate the end of L
  while left < right
  {
    while left < right and S[left] < pivot
      left ++;         // insert into L
    while left < right and S[right] ≥ pivot
      right --;        // insert into G
    if left < right
      Swap S[left] and S[right];
  }
  Swap S[left] and pivot;
  return left;
}
```



In-Place Partitioning (2/3)



- Perform the partition using two indices to split S into L and $E \cup G$ (a similar method can split $E \cup G$ into E and G).




- Repeat until *left* and *right* meet:
 - Scan left to the right until finding an element \geq pivot.
 - Scan right to the left until finding an element $<$ pivot.
 - Swap elements at indices *left* and *right*



In-Place Partitioning (2/3)




- Swap the element at *left* and pivot



Summary of Sorting Algorithms

Algorithm	Time	Notes
selection-sort	$O(n^2)$	<ul style="list-style-type: none">■ in-place■ slow (good for small inputs)
insertion-sort	$O(n^2)$	<ul style="list-style-type: none">■ in-place■ slow (good for small inputs)
quick-sort	$O(n \log n)$ expected	<ul style="list-style-type: none">■ in-place, randomized■ fastest (good for large inputs)
heap-sort	$O(n \log n)$	<ul style="list-style-type: none">■ in-place■ fast (good for large inputs)
merge-sort	$O(n \log n)$	<ul style="list-style-type: none">■ sequential data access■ fast (good for huge inputs)

Bucket-Sort and Radix-Sort



Bucket-Sort



- Let be S be a sequence of n (key, value) entries with keys in the range $[0, N - 1]$
- Bucket-sort uses the keys as indices into an auxiliary array B of sequences (buckets)

Phase 1: Empty sequence S by moving each entry (k, o) into its bucket $B[k]$

Phase 2: For $i = 0, \dots, N - 1$, move the entries of bucket $B[i]$ to the end of sequence S

- Analysis:
 - Phase 1 takes $O(n)$ time
 - Phase 2 takes $O(n + N)$ time
- Bucket-sort takes $O(n + N)$ time

Algorithm *bucketSort*(S, N)

Input sequence S of (key, value) items with keys in the range $[0, N - 1]$
Output sequence S sorted by increasing keys

{ B = array of N empty sequences;
while (\neg *isEmpty*(S)

{ $f = \text{first}(S)$;
 $(k, o) = \text{remove}(S, f)$;
 $\text{insertLast}(B[k], (k, o))$; }

for ($i = 0$; $i++$; $i \leq N - 1$)


while (\neg *isEmpty*($B[i]$))
 { $f = \text{first}(B[i])$;
 $(k, o) = \text{remove}(B[i], f)$;
 $\text{insertLast}(S, (k, o))$; }

}


Example




- Key range [0, 9]



Phase 1



Phase 2



Properties and Extensions




- Key-type Property
 - The keys are used as indices into an array and cannot be arbitrary objects
- **Stable** Sort Property
 - The relative order of any two items with the same key is preserved after the execution of the algorithm

Extensions

- Integer keys in the range $[a, b]$
 - Put entry (k, o) into bucket $B[k - a]$
- String keys from a set D of possible strings, where D has constant size (e.g., names of the 50 U.S. states)
 - Sort D and compute the rank $r(k)$ of each string k of D in the sorted sequence
 - Put entry (k, o) into bucket $B[r(k)]$

Lexicographic Order



- A d -tuple is a sequence of d keys (k_1, k_2, \dots, k_d) , where key k_i is said to be the i -th dimension of the tuple
- Example:
 - The Cartesian coordinates of a point in space are a 3-tuple
- The lexicographic order of two d -tuples is recursively defined as follows

$$(x_1, x_2, \dots, x_d) < (y_1, y_2, \dots, y_d)$$



$$x_1 < y_1 \vee x_1 = y_1 \wedge (x_2, \dots, x_d) < (y_2, \dots, y_d)$$

i.e., the tuples are compared by the first dimension, then by the second dimension, etc.

Lexicographic-Sort

- Let *stableSort*(*S*) be a stable sorting algorithm
- Lexicographic-sort sorts a sequence of *d*-tuples in lexicographic order by executing *d* times algorithm *stableSort*, one per dimension, from least significant element to most significant element
- Lexicographic-sort runs in $O(dT(n))$ time, where $T(n)$ is the running time of *stableSort*

Algorithm *lexicographicSort*(*S*)

Input sequence *S* of *d*-tuples

Output sequence *S* sorted in lexicographic order

```
{ for ( i = d; i >= 1; i - -; )  
    stableSort(S, i);  
}
```

Example:


(7,4,6) (5,1,5) (2,4,6) (2, 1, 4) (3, 2, 4)

(2, 1, 4) (3, 2, 4) (5,1,5) (7,4,6) (2,4,6)

(2, 1, 4) (5,1,5) (3, 2, 4) (7,4,6) (2,4,6)

(2, 1, 4) (2,4,6) (3, 2, 4) (5,1,5) (7,4,6)

Radix-Sort



- Radix-sort is a specialization of lexicographic-sort that uses bucket-sort as the stable sorting algorithm in each dimension
- Radix-sort is applicable to tuples where the keys in each dimension i are integers in the range $[0, N - 1]$
- Radix-sort runs in time $O(d(n + N))$

Algorithm *radixSort*(S, N)

Input sequence S of d -tuples such that $(0, \dots, 0) \leq (x_1, \dots, x_d)$ and $(x_1, \dots, x_d) \leq (N - 1, \dots, N - 1)$ for each tuple (x_1, \dots, x_d) in S

Output sequence S sorted in lexicographic order

```
{ for (  $i = d$ ;  $i \geq 1$ ;  $i--$  )  
    bucketSort( $S, N$ );  
}
```

Radix-Sort for Binary Numbers



- Consider a sequence of n b -bit integers
$$x = x_{b-1} \dots x_1 x_0$$
- We represent each element as a b -tuple of integers in the range $[0, 1]$ and apply radix-sort with $N = 2$
- This application of the radix-sort algorithm runs in $O(bn)$ time
- For example, we can sort a sequence of 32-bit integers in linear time

Algorithm *binaryRadixSort(S)*

Input sequence S of b -bit integers

Output sequence S sorted
replace each element x of S with the item $(0, x)$


{ **for** ($i = 0; i \leq b-1; i++$)

{ replace the key k of each item (k, x) of S with bit x_i of x ;

bucketSort(S, 2);


}

}




Example

- Sorting a sequence of 4-bit integers




Sorting Lower Bound



Comparison-Based Sorting




- Many sorting algorithms are comparison based.
 - They sort by making comparisons between pairs of objects
 - Examples: bubble-sort, selection-sort, insertion-sort, heap-sort, merge-sort, quick-sort, ...
- Let us therefore derive a lower bound on the running time of any algorithm that uses comparisons to sort n elements, x_1, x_2, \dots, x_n .




Counting Comparisons

- Let us just count comparisons then.
- Each possible run of the algorithm corresponds to a root-to-leaf path in a **decision tree**




Decision Tree Height

- The height of this decision tree is a lower bound on the running time
- Every possible input permutation must lead to a separate leaf output.
 - If not, some input ...4...5... would have same output ordering as ...5...4..., which would be wrong.
- Since there are $n! = 1 * 2 * \dots * n$ leaves, the height is at least $\log(n!)$



The Lower Bound



- Any comparison-based sorting algorithm takes at least $\log(n!)$ time
- Therefore, any such algorithm takes time at least

$$\log(n!) \geq \log\left(\frac{n}{2}\right)^{\frac{n}{2}} = (n/2) \log(n/2).$$

- That is, any comparison-based sorting algorithm must run in $\Omega(n \log n)$ time.



Summary

- Merge sort
- Quick sort
- Lexicographic sort
- Bucket sort
- Radix sort

Suggested reading:

Sedgewick, Chapters 7, 8, 10